

Chung kết Cuộc thi  
**OLYMPIC TRÍ TUỆ NHÂN TẠO - OAI HCMC 2025**  
(Olympiad in Artificial Intelligence at Ho Chi Minh City 2025)

**TP. Hồ Chí Minh, ngày 25 tháng 04 năm 2025**





VIETNAM NATIONAL UNIVERSITY HO CHI MINH CITY  
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY



HCMUT EE  
Machine Learning & IoT Lab

# MUSHROOM CLASSIFICATION

**Presenter: EEIoT\_newbie**



# EEIoT\_newbie



**VĂN THỊNH**



**DUY TÂN**



**ANH QUÂN**



**MINH HUY**



# Outline

---

- ◆ 1. Introduction
- ◆ 2. Problem Statement
- ◆ 3. Method
- ◆ 4. Experimental Results

# 1. Introduction

# 1. Introduction

---

**Topic:** Classification of dried mushroom images.

**Objective:** Build a machine learning model to classify images into three categories:

- **Shiitake Mushroom** (label 0)
- **Wood Ear Mushroom** (label 1)
- **Snow Fungus** (label 2)

**Dataset:**

- Color images, **32x32 pixels** resolution.
- **Training set:** 1050 images.
- **Testing set:** 450 images.

## 2. Problem Statement

## 2. Problem Statement

### *Input:*

- **Training:**
  - 500 RGB color images, size **32x32 pixels**, equally distributed among: Snow Fungus (167), Wood Ear Mushroom (166), Shiitake Mushroom (167).
  - All images have uniform dimensions, 1:1 aspect ratio, file size between **0.80 – 1.04 KB**.
- **Inference:**
  - 450 unlabeled test images with the same resolution and RGB format.
  - Statistical differences in **brightness** and **contrast** across classes:
    - Shiitake: brightest (114.09), lowest contrast (41.13)
    - Snow Fungus: moderate brightness (102.14), medium contrast (52.84)
    - Wood Ear: darkest (101.63), highest contrast (59.09)

*Output:* Predict the correct label (class) for each test image:

- **0:** Shiitake Mushroom
- **1:** Wood Ear Mushroom
- **2:** Snow Fungus



# 3. Method

# 3. Method

---

## ❖ Calculating Mean and Standard Deviation for Data Normalization

### Why Normalize?

Helps deep learning models **train faster and more stably**.

Ensures all input features (pixels) are on a similar scale.

Typically transforms image pixel values so each channel (RGB) has **mean = 0** and **standard deviation = 1**.

Prevents issues like exploding or vanishing gradients.

# 3. Method

## Calculating Mean and Standard Deviation for Data Normalization

### Theoretical Formulas:

- Mean ( $\mu$ ) per channel:

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i$$


- Standard Deviation ( $\sigma$ ) per channel:

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$$

Where  $x_i$  is a pixel value, and  $N$  is the total number of pixels across all images per channel.

# 3. Method

- ❖ RandomHorizontalFlip
- ❖ RandomVerticalFlip
- ❖ ResizeCrop



### 3. Method

#### ❖ ColorJitter

Original Image




Transformed Image



### 3. Method


#### ❖ RandomErasing (Cutout)




# 3. Method

## ❖ Elastic Transform

Original Image

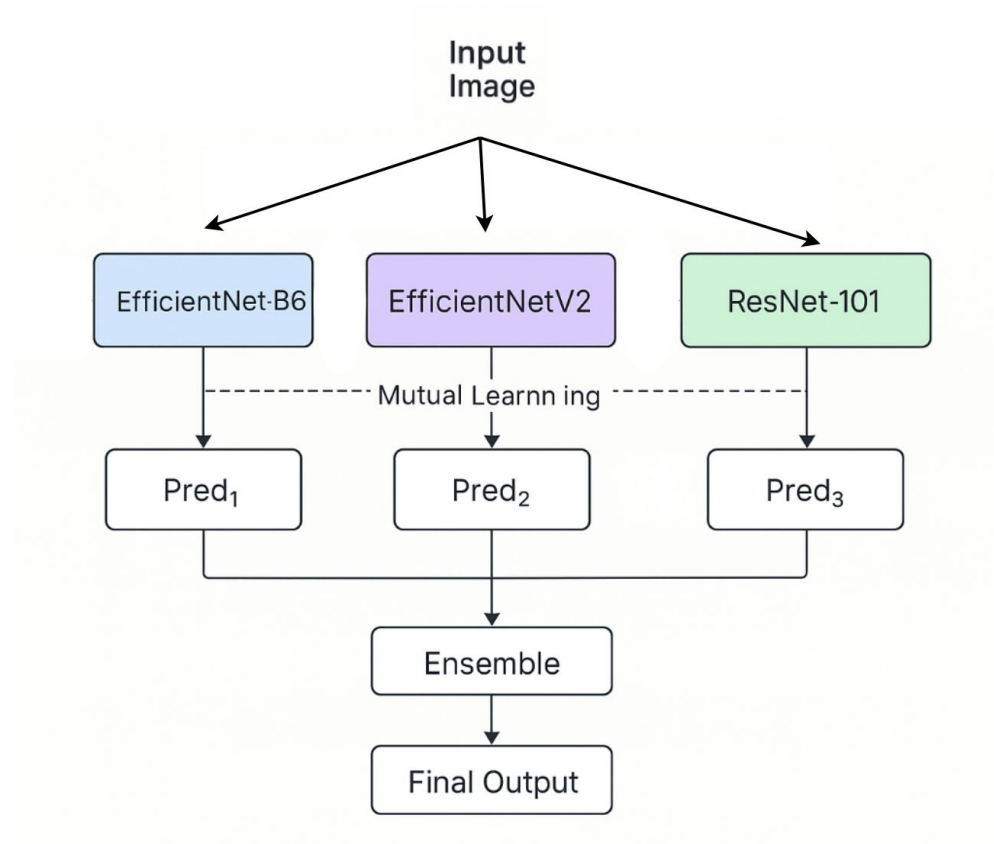


Transformed Image



# 3. Method

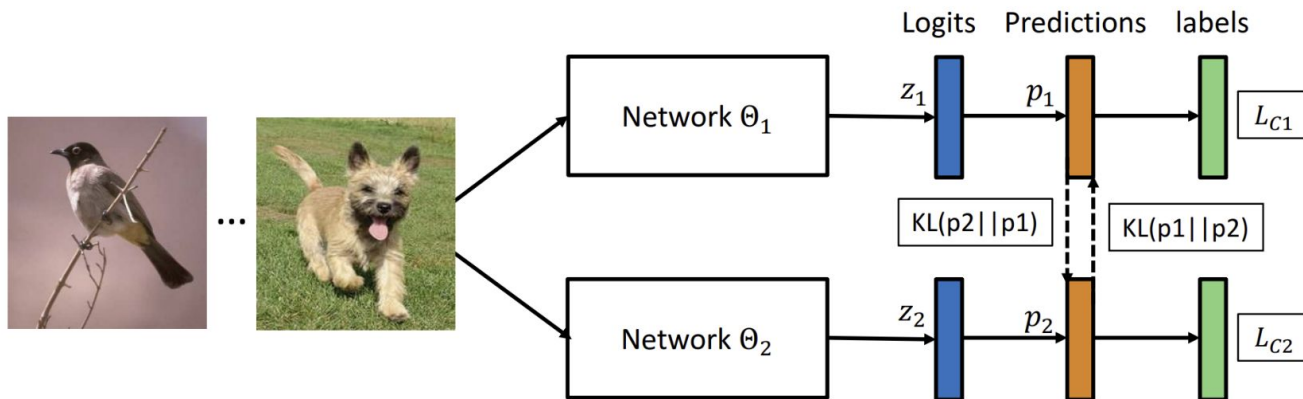
## ❖ Overview





### 3. Method

#### ❖ Deep Mutual Learning

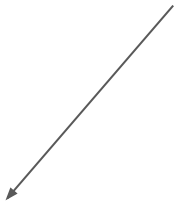


# 3. Method

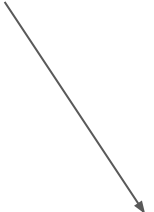
## ◆ Deep Mutual Learning

4. The total loss for model  $\Theta_i$  with  $K$  networks:

$$L_{\Theta_i} = L_{C_i} + \frac{1}{K-1} \sum_{k=1, k \neq i}^K D_{KL}(p_k \parallel p_i)$$


$$L_{C_i} = - \sum_{i=1}^N \sum_{m=1}^M I(y_i, m) \log(p_i^m(x_i))$$

Cross-Entropy


$$D_{KL}(p_k \parallel p_i) = \sum_{i=1}^N \sum_{m=1}^M p_k^m(x_i) \log \frac{p_k^m(x_i)}{p_i^m(x_i)}$$

Kullback-Leibler (KL) divergence

# 3. Method

## ◆ Ensemble Learning with Unweighted Averaging

$$p = \frac{1}{K} \sum_{k=1}^K p_k$$

$K$  : the number of models

$p_k$  : the probabilities are calculated from model  $\mathbf{k}$

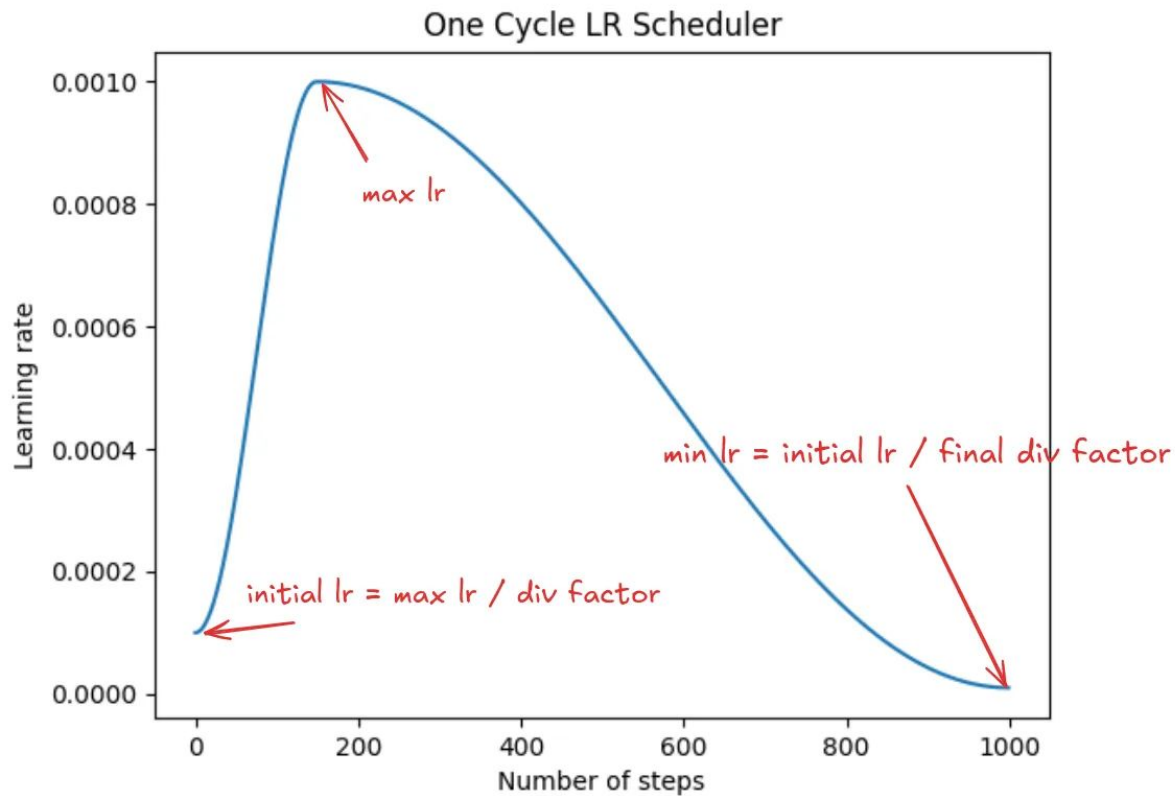
## 4. Experimental Results

## 4. Experimental Results

- **Image size : 224 x 224**
- **Batch size : 32**
- **LossFunction : CrossEntropyLoss(label\_smoothing=0.1)**
- **Optimizer : AdamW(lr=0.001, weight\_decay=0.01)**
- **EarlyStopping : finish training when valid\_acc does not improve (3 patience)**
- **clip\_grad\_norm : normalize gradient**

## 4. Experimental Result

### ◆ OneCycleLR



## 4. Experimental Result

### ❖ GradScaler - Mixed Precision Training

Use float16 instead float32 gradient

- Scale loss : multiplying gradient to *scale factor*-> avoid vanishing gradient
- Unscale loss
- Adjust scale factor flexibly

=> Faster training

=> Optimize GPU utilization

## 4. Experimental Result

### ◆ Public Test

Method	Num_epochs	Optimizer	DML	Ensemble	Test Accuracy
Resnet-34	30	SGD	✗	✗	0.76
ViT-L16	20	SGD	✗	✗	0.83
GoogleNet	20	Adam	✗	✗	0.85
MobileNet	15	Adam	✗	✗	0.85
EfficientNetB6-CBAM	20	SGD	✗	✗	0.91
EfficientNetB6	17	SGD	✗	✗	0.96
MobileNet + EfficientNetB6 + VGG16	50	Adam	✗	✓	0.95
EfficientNetB6* + EfficientNetV2 + Resnet101	15	AdamW	✓	✗	0.95
EfficientNetB6 + EfficientNetV2 + Resnet101	15	AdamW	✓	✓	<b>0.98</b>



## 4. Experimental Result

### ❖ Private Test

Method	Num_epochs	Optimizer	DML	Ensemble	Test Accuracy
EfficientNetB6 + EfficientNetV2 + Resnet101	15	AdamW	✓	✓	<b>0.96</b>

			Accuracy	F1-Score
4	OAI069	EEIoT_newbie	96.17%	96.10%

## 4. Experimental Result

### ❖ Private Test (Final)

Method	Num_epochs	Optimizer	DML	Ensemble	Test Accuracy
EfficientNetB6 + EfficientNetV2 + Resnet101	15	AdamW	✓	✓	<b>0.96</b>